

Gefördert durch:



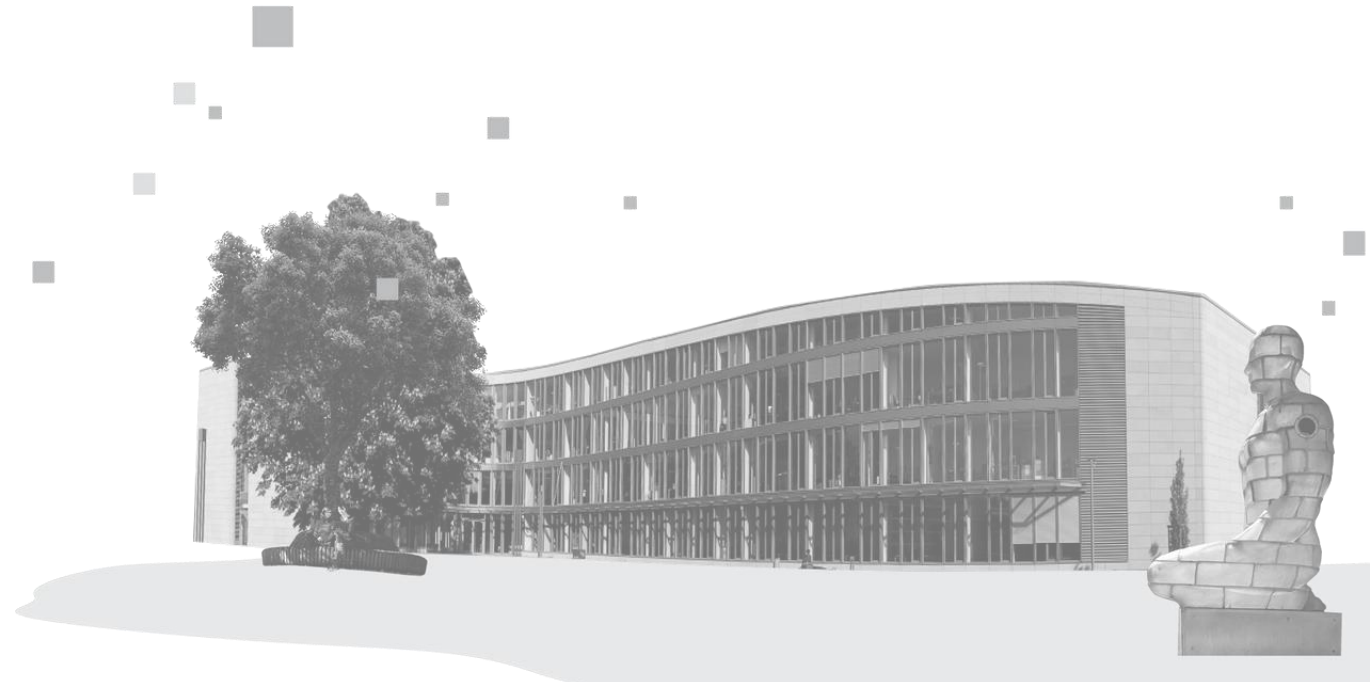
# Time Series Forecasting

## 3.3 Convolutional Networks

Mario Tormo Romero

**Design IT.  
Create Knowledge.**

[www.hpi.de](http://www.hpi.de)



# What we'll cover in this video




- Why CNNs are useful for time series data
- Key concepts:
  - 1D and dilated convolutions
  - Residual connections and pooling
- Introduction to Temporal Convolutional Networks (TCNs)
- Practical example: Building a CNN model for forecasting

# Why CNNs Are Useful for Time Series Forecasting




- Pattern detection through sliding filters
- Parallel processing of sequences — faster training compared to RNNs
- Reduced vanishing/exploding gradient issues in deeper models
- Efficient capture of local temporal patterns
- Dilated convolutions expand receptive field for long-range dependencies




# What Are Convolutional Neural Networks?

- Designed to automatically and adaptively learn spatial and temporal patterns through convolutional layers.
- How They Work:
  - Apply filters (kernels) that slide over input data to extract local features.
  - Each filter detects specific patterns (like edges in images or temporal trends in sequences).
  - Multiple filters produce different feature maps, capturing diverse characteristics.
  - Stacking convolutional layers enables learning of increasingly complex patterns.
- Key Characteristics:
  - **Local Connectivity:** Focus on small neighborhoods in data at a time.
  - **Parameter Sharing:** Same filter used across the entire input reduces the number of parameters.
  - **Hierarchical Feature Learning:** Lower layers learn simple features, higher layers learn complex abstractions.



# Core Components of CNNs for Time Series Forecasting

- **1D Convolutions**
  - Filters slide along the time axis
  - Extract local temporal patterns
- **Dilated Convolutions**
  - Skip input points to expand receptive field
  - Capture longer-range dependencies without deeper networks
- **Residual Connections**
  - Bypass pathways to help train deeper CNNs
  - Prevent vanishing gradients and information loss
- **Pooling Layers**
  - Downsample the feature maps
  - Reduce noise and computational cost



# Core Components of CNNs for Time Series Forecasting

- **1D Convolutions**

- Filters slide along the time axis
- Extract local temporal patterns

- **Dilated Convolutions**


- Skip input points to expand receptive field
- Capture longer-range dependencies without deeper networks

- **Residual Connections**

- Bypass pathways to help train deeper CNNs
- Prevent vanishing gradients and information loss


- **Pooling Layers**

- Downsample the feature maps
- Reduce noise and computational cost



# Core Components of CNNs for Time Series Forecasting

- **1D Convolutions**
  - Filters slide along the time axis
  - Extract local temporal patterns
- **Dilated Convolutions**
  - Skip input points to expand receptive field
  - Capture longer-range dependencies without deeper networks
- **Residual Connections**
  - Bypass pathways to help train deeper CNNs
  - Prevent vanishing gradients and information loss
- **Pooling Layers**
  - Downsample the feature maps
  - Reduce noise and computational cost



# Core Components of CNNs for Time Series Forecasting

- **1D Convolutions**

- Filters slide along the time axis
- Extract local temporal patterns

- **Dilated Convolutions**

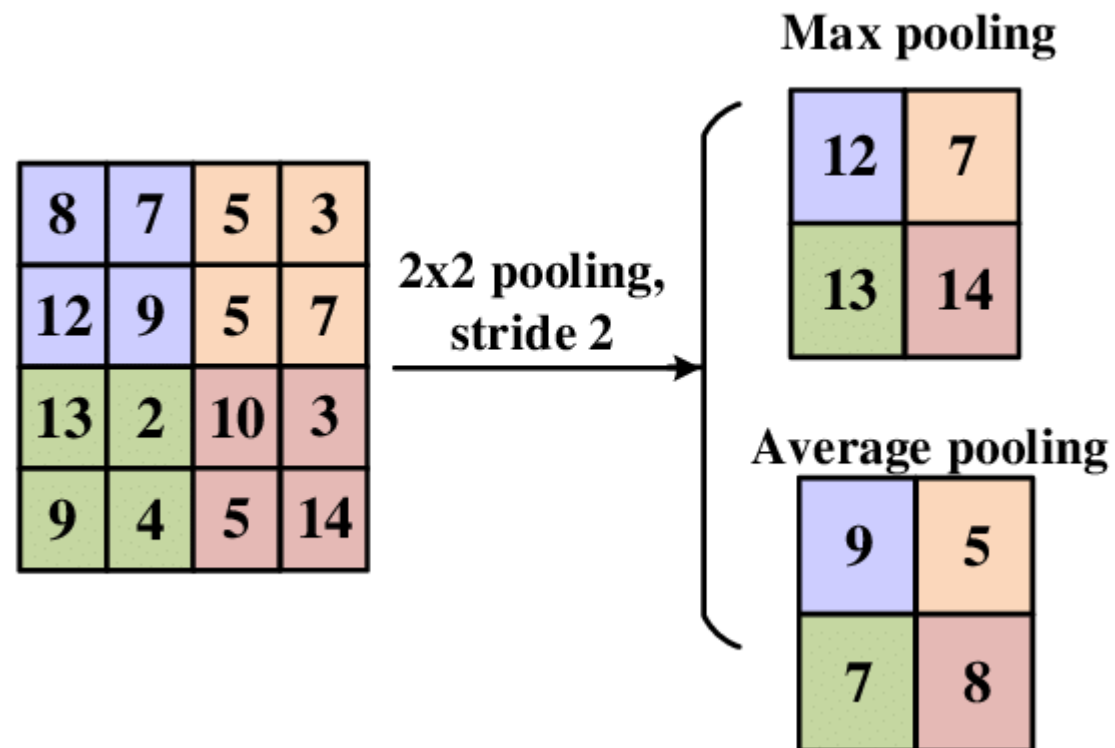
- Skip input points to expand receptive field
- Capture longer-range dependencies without deeper networks

- **Residual Connections**

- Bypass pathways to help train deeper CNNs
- Prevent vanishing gradients and information loss

- **Pooling Layers**


- Downsample the feature maps
- Reduce noise and computational cost





# Temporal Convolutional Networks (TCNs)

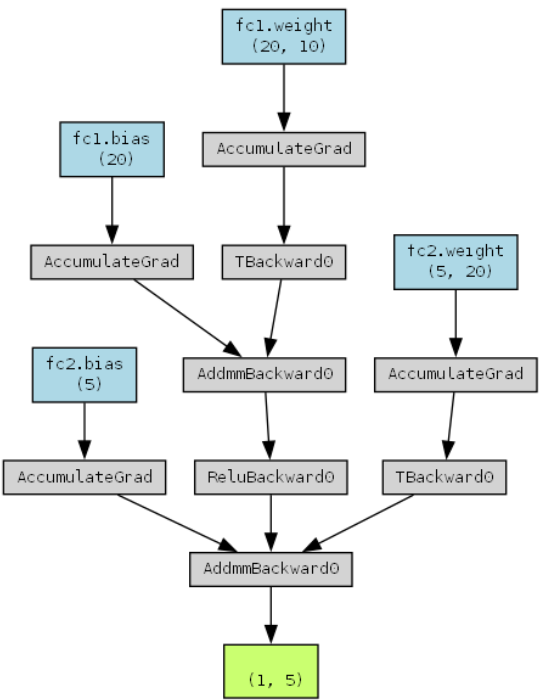
- CNN architecture specifically designed for sequence modeling
- Uses **causal convolutions** to ensure no future information leaks into predictions
- Employs **dilated convolutions** to capture long-range dependencies
- Incorporates **residual connections** for stable and deep networks
- **Parallelizable**: faster training than RNNs/LSTMs
- Avoid vanishing gradient problems common in recurrent models
- Often achieve state-of-the-art performance on sequence tasks



# Practical Example — Implementing a CNN for Time Series Forecasting

- Dataset: Univariate time series from **Berlin Daily Temperature Dataset**
- Model Architecture:
  - Input Layer** (samples, time steps, features)
  - 1D Convolutional Layer** (extract local temporal patterns)
  - Max Pooling Layer** (reduce dimensionality)
  - Fully Connected Layer** (map to forecast output)
- Framework: **Python + PyTorch**
- Goal:** Predict the next value from past observations

```
CNNForecast(
  (conv1): Conv1d(1, 16, kernel_size=(3,), stride=(1,))
  (relu): ReLU()
  (pool): MaxPool1d(kernel_size=2, stride=2, padding=0,
                    dilation=1, ceil_mode=False)
  (fc): Linear(in_features=224, out_features=1, bias=True)
)
```



# What we've learnt



- CNNs efficiently capture local temporal patterns in time series data.
- 1D and dilated convolutions expand receptive fields without extra parameters.
- Temporal Convolutional Networks (TCNs) use causal convolutions and residual blocks for better long-range dependency modeling.
- CNNs enable faster, parallel training compared to RNNs.
- Practical CNN models for forecasting include convolutional, pooling, and dense layers implemented in frameworks like PyTorch.